

Lab: Pthreads Synchronization

Goal

In this lab, you will learn the basics of multi-threaded programming, and synchronizing multiple threads using locks and condition variables. You will use the `pthread` thread API in this assignment.

Before you begin

Please familiarize yourself with the `pthread` API thoroughly.

Part A: Master-Worker Thread Pool

In this part of the lab, you will implement a simple master and worker pool, a pattern that occurs in many real life applications. You are given a program that spawns one master/producer thread. The master produces numbers continuously. You must complete this program in the following manner: you must store the produced numbers in a buffer, and spawn a pool of worker threads which will consume (i.e., print) these numbers to the screen. This simple program is an example of a master-worker thread pool pattern that is commonly found in several real-life application servers. For example, in the commonly used multi-threaded architecture for web servers, the server has one master thread and a pool of worker threads. When a new connection arrives from a web client, the master hands the request over to one of the workers. The worker then reads the web request from the network socket and writes a response back to the client. Your simple master-worker program is similar in structure to such applications, albeit with much simpler request processing logic at the application layer.

You are given a skeleton program `master-worker-skeleton.c`. This program takes 3 command line arguments: how many numbers “produce” (M), the maximum size of the buffer in which the produced numbers should be stored (N), and the number of worker threads to spawn to consume these numbers (W). The skeleton code spawns a master thread, that produces the specified number of integers from 0 to $M - 1$ and exits. The main program waits for this thread to join and terminates.

The master and worker threads must synchronize using a shared buffer. The shared buffer can either be implemented as a fixed size array of N integers, or a dynamic sized linked list. If you are using the linked list option, the linked list must start with zero elements. The producer thread must dynamically add elements to the list when producing integers, and the worker threads must delete elements from the list when consuming. Note that you must implement the linked list operations yourself, and must not use any library implementation of lists. You must dynamically allocate and free memory for the elements of the list. All insert/delete operations on the linked list must be implemented with proper locking to avoid

race conditions. The maximum size of the linked list must still be bounded by the maximum buffer size N .

Copy the skeleton code provided to you either into `master-worker-array.c` or `master-worker-ll.c`, depending on your design choice. Now, add the following functionality to the skeleton code. First, you must suitably initialize the shared buffer. The master thread places produced integers into this buffer. The worker threads must remove integers from this buffer to consume and print to screen. The master thread must wait on a condition variable if the buffer is full, and the worker threads must similarly wait if the buffer is empty. The master and workers must coordinate to ensure that every request is processed exactly once and only once by some worker. While you need to ensure that all W workers are involved in consuming the integers, it is not necessary to ensure perfect load balancing between the workers. Once all M integers (from 0 to $M - 1$) have been produced and consumed, all threads must exit. The main thread must call `pthread_join` on the master and worker threads, and terminate itself once all threads have joined.

The master thread must call the function `print_produced` when it produces an integer into the buffer, and the worker threads must call the function `print_consumed` when it removes an integer from the buffer to consume. If your code is written correctly, every integer from 0 to $M - 1$ will be printed exactly once by the master producer thread, and exactly once by the worker consumer threads. We have provided you with a simple autograding script that checks this above invariant on the output produced by your program. **Please do not modify these print functions, as their output will be parsed by our grading script.**

Part B: Reader-Writer Locks

Consider an application where multiple threads of a process wish to read and write data shared between them. Some threads only want to read (let's call them "readers"), while others want to update the shared data ("writers"). In this scenario, it is perfectly safe for multiple readers to concurrently access the shared data, as long as no other writer is updating it. However, a writer must still require mutual exclusion, and must not access the data concurrently with any other thread, whether a reader or a writer. A reader-writer lock is a special kind of a lock, where the acquiring thread can specify whether it wishes to acquire the lock for reading or writing. That is, this lock will expose two separate locking functions, say, `ReaderLock()` and `WriterLock()`, and analogous unlock functions. If a lock is acquired for reading, other threads that wish to read may also be permitted to acquire the lock at the same time, leading to improved concurrency over traditional locks.

There are two flavors of the reader-writer lock, which we will illustrate with an example. Suppose a reader thread $R1$ has acquired a reader-writer lock for reading. While $R1$ holds this lock, a writer thread W and another reader thread $R2$ have both requested the lock. Now, it is fine to allow $R2$ also to simultaneously acquire the lock with $R1$, because both are only reading shared data. However, allowing $R2$ to acquire the lock may prolong the waiting time of the writer thread W , because W has to now wait for both $R1$ and $R2$ to release the lock. So, whether we wish to permit more readers to acquire the lock when a writer is waiting is a design decision in the implementation of the lock. When a reader-writer lock is implemented with *reader preference*, additional readers are allowed to concurrently hold the lock with previous readers, even if a writer thread is waiting for the lock. In contrast, when a reader-writer lock is implemented with *writer preference*, additional readers are not granted the lock when a writer is already waiting. That is, we do not prolong the waiting time of a writer any more than required.

In this part of the lab, you must implement both flavors of the reader-writer lock. You must complete the definition of the structure that captures the reader-writer lock in `rwlock.h`. The following functions

to be supported by this lock are also defined in the header file:

- The function `InitializeReadWriteLock()` must initialize the lock suitably.
- The function `ReaderLock()` is invoked by a reader thread before entering a read-only critical section, and the function `ReaderUnlock()` is invoked by the reader when exiting the critical section.
- The function `WriterLock()` is invoked by a writer thread before entering a critical section with shared data updates, and the function `WriterUnlock()` is invoked by the writer when exiting the critical section.

You must write the code to implement these functions. You must write code that implements reader-writer locks with reader preference in `rwlock-reader-pref.cpp`. Similarly, you must implement reader-writer locks with writer preference in `rwlock-writer-pref.cpp`. Note that both implementations of the lock must share the same header file, including the definition of the reader-writer lock structure. Therefore, your lock structure may have some fields that are only used in one version of the code and not the other.

As part of the autograding scripts given to you, you have been provided with a C++ program `tester.cpp` to test your reader-writer lock implementation. You can compile and link this tester with both flavors of the lock separately to create two different executables of the tester, as follows.

```
$g++ tester.cpp rwlock-reader-pref.cpp -o rwlock-reader-pref -lpthread
$g++ tester.cpp rwlock-writer-pref.cpp -o rwlock-writer-pref -lpthread
```

The tester programs takes two command line arguments, say R and W . The program then spawns R reader threads, followed by W writer threads, followed by R *additional* reader threads. Each thread, upon creation, tries to acquire the same reader-writer lock as a reader or writer (as the case may be), holds the lock for a long period of time, and finally releases the lock. The tester program prints a message to the screen when it acquires and releases the lock, and the relative ordering of these print messages will help us verify the correctness of your implementation. By invoking this tester with different values of M and N , one can test the reader-writer reasonably thoroughly. Of course, you are encouraged to write your own test programs that use the reader-writer lock as well.

We have also provided you with an autograder script `autograder.sh`, which does the following. This script compiles the tester program with both flavors of the reader-writer lock, runs the two tester executables for different values of M and N , and stores the print output. It then runs the autograder scripts `autograder-reader-pref.py` and `autograder-writer-pref.py` on these outputs. These Python scripts parse the output printed from the tester program, and ensure that all the threads have successfully acquired and released locks in the expected order.

Submission instructions

- For part A, you must submit ONE of these two files: `master-worker-array.c` OR `master-worker-ll.c`. For part B, you must submit all the files `rwlock.h`, `rwlock-reader-pref.cpp`, and `rwlock-writer-pref.cpp`.
- Place these files and any other files you wish to submit in your submission directory, with the directory name being your roll number (say, 12345678).

- Tar and gzip the directory using the command `tar -zcvf 12345678.tar.gz 12345678` to produce a single compressed file of your submission directory. Submit this tar gzipped file on Moodle.

Grading

We will check correctness of your code using the test scripts provided (with additional test cases beyond those provided to you). In addition to using the test scripts, we will also read through your code to ensure that you have adhered to the problem specification in your implementation.